


SHAP by hand — worksheet

OPIM 5512 · Data Science Using Python · Dr. Dave Wanik — name: _____

SHAP explains **one** prediction: how much did each feature push it **above or below what we'd normally expect**? The "fair" split comes from **Shapley values** — a feature's *average marginal contribution* over every coalition it could join. Here you compute the whole thing by hand from a **printed decision tree**; you'll then recognize the exact same steps (on a real California-housing tree) in the notebook `2a_SHAP_ShallowTree_CAHousing`.

The model — one shallow tree

A tree scores a home (median value, in **units of 100,000 USD**) from three features: **I** = income (in units of 10,000 USD), **R** = rooms, **L** = latitude.



The value function — how we score a coalition

To score a coalition S (the features we treat as **"known"**), we set the features in S to **this home's** values and the rest to a **baseline** row, then read the tree's leaf. (The notebook averages over a whole *background sample*; here one baseline row keeps the arithmetic by hand.)

feature	baseline row	this home x
I (income, 10k USD)	3	8
R (rooms)	3	6
L (latitude)	34	38

So $v(\emptyset)$ = trace the tree with the **all-baseline** row (nothing known), and $v(\{I, R, L\})$ = trace it with the **all-home** row (everything known) = the prediction.

The Shapley value (coalition form), for $p = 3$ features:

$$\phi_i = \sum_{S \subseteq N \setminus \{i\}} w(|S|) [v(S \cup \{i\}) - v(S)], \quad w(|S|) = \frac{|S|! (p - |S| - 1)!}{p!}.$$

How to read that formula

One sentence: "for every team feature i could join, measure i 's boost, weight it, and add them up."

Symbol by symbol:

symbol	what it means
N	all the features: $\{I, R, L\}$
$N \setminus \{i\}$	" N without i " — every feature <i>except</i> i (the $\setminus$ is set-minus)
$S \subseteq N \setminus \{i\}$	a team of the <i>other</i> features; the $\sum$ loops over all of them
$S \cup \{i\}$	union — the team S with i added in (e.g. $\{R\} \cup \{I\} = \{I, R\}$)
$v(S \cup \{i\}) - v(S)$	value with i minus without i = i 's boost to that team
$w(S)$	the weight — depends only on the team's size

Where the union $\cup$ comes in: it's just the "+ i joins the team" step. You trace the tree once with team S known ($v(S)$, so i is on baseline), then **again** with S **and** i known ($v(S \cup \{i\})$) — the difference is what i added.

The weights are constant. $w(|S|)$ depends only on the team **size** (and p) — **not** on the house. For $p = 3$ they are **always** $\frac{1}{3}, \frac{1}{6}, \frac{1}{6}, \frac{1}{3}$, reused for every prediction. Only the $v(S)$ numbers change from row to row.

Tasks

1. **Trace the tree** to fill in the value of all 8 coalitions. For each, set the listed features to the home's value and the rest to baseline, then follow the tree to a leaf.

S	$\{\}$	$\{I\}$	$\{R\}$	$\{L\}$	$\{I, R\}$	$\{I, L\}$	$\{R, L\}$	$\{I, R, L\}$
$v(S)$								

2. **Weights.** Compute $w(|S|)$ for $|S| = 0, 1, 2$ (with $p = 3$).
3. ϕ_I . List the four terms $w \cdot [v(S \cup \{I\}) - v(S)]$ for $S = \{\}, \{R\}, \{L\}, \{R, L\}$ and add them.
4. ϕ_R **and** ϕ_L . Same for rooms and latitude.
5. **Efficiency (local accuracy).** Check that **base value** + $\phi_I + \phi_R + \phi_L =$ **prediction**.
6. **Interpret the interactions.** Explain *from the tree* why $v(\{I, R\}) = v(\{I\})$ (does the tree ever look at R once income is high?), and why $v(\{L\}) = v(\emptyset)$.