


Lab 3: Auditing Your ETL Pipeline

Does your pipeline see what *you* see?

OPIM 5512 · Module 3 · Solo · Colab · your own Google Cloud bucket

Tonight you measure your extractor — then fix it with one regular expression.

Your pipeline has been busy. Is it any good?



Part	Time	You do
1 · Get it running	30 min	5 green checks, then Force run scraper → extractor → materialize
2 · Search and see	25 min	your cars land in Colab's Files panel · search them · open 3 live ads
3 · Validate the extractor	30 min	a truth table for 5 cars → accuracy per field
4 · Fix it three ways	30 min	plain Python (given) · your one-line regex · a typo-fixer → re-grade

Pipeline not ready? One switch: `USE_BACKUP_DATA = True` loads Dr. Wanik's Hartford cars. Nobody sits out.

Get it running (if it isn't yet)

- 1 • Final GCP Guide F26 (HuskyCT, Week 3.1), Steps 0–10 → **5 green checks** in your fork's Actions tab
- 2 • Cloud Scheduler: **tick the checkbox first**, then **Force run** — one at a time, wait for **Success**:
① `craigslist-scraper-hourly` ② `extractor-per-listing-hourly` ③ `materialize-master-hourly`
- 3 • Cloud Storage → your bucket: `scrapes/` and `structured/` are there
- 4 • Already done? Skip to Part 2 and help a neighbor.

Stuck? Ask Claude first: paste the step, the command, and the FULL error. Then raise your hand.

A map of your bucket

```
gs://your-project-id/  
├ scrapes/  
│   └ 20260928155246/           # one folder per scraper run (UTC)  
│       ├── eH4UEPtCwKz2Vox1upEYCw.txt # the raw ad text  
│       └ index.csv              # which URL each file came from  
├ state/seen_post_ids.txt       # how it never grabs a car twice  
└ structured/  
    ├── run_id=20260928155246/jsonl/ # the regex's answers, per ad  
    ├── datasets/listings_master.csv  # ← tonight's table  
    └ preds/preds_master.csv         # predictions (once you have 2 days)
```

Raw → structured → one table. That's ETL. Tonight you check the T.

Search and see

IN YOUR NOTEBOOK

```
SEARCH_WORD = "toyota" # try sunroof, 4wd...

for i in range(len(cars)):
    row = cars.iloc[i]
    text = read_ad_text(row["source_txt"])
    if SEARCH_WORD.lower() in text.lower():
        matches.append(row)

2 of 29 ads mention 'toyota'
```

FILES PANEL → MY_CARS/01_2013_8500_....TXT

```
2013 Hyunda Sonata GLS Limited for sale by owner
$8,500
(East Hartford)
image 1 of 13
2013
Hyunda Sonata GLS Limited
condition:
like new
odometer:
40,000
```

Double-click any car on the left: the .txt is the raw ad, the .json is what the extractor pulled out, and every car keeps its link to the live ad.

Validate the extractor: ground truth

Real Hartford cars. The extractor nails the numbers. Make and model? **Vernon Rockville** is a town. **Hyunda** is the seller's typo:

post_id	price	year	make	model	mileage
Sonata	8500 ✓	2013 ✓	Hyunda ✗	Sonata ✓	40000 ✓
Accord	3800 ✓	2011 ✓	Honda ✓	Accord ✓	198000 ✓
Elantra	4995 ✓	2013 ✓	Contact ✗	Information ✗	88000 ✓
Odyssey	1800 ✓	2007 ✓	Vernon ✗	Rockville ✗	213000 ✓
accuracy (5 cars)	100%	100%	20%	40%	100%

Truth-table rule: write the REAL make. If the seller typed Hyunda, the truth is hyundai. Lowercase; model = first word.

THE MYSTERY

What is the make/model regex really grabbing?

THE EXTRACTOR'S RULE (MODULE 3.2)

```
MAKE_MODEL_RE = re.compile(
    r"\b([A-Z][a-z]+)\s+([A-Z][A-Za-z0-9]+)")

# in words: the FIRST place where one
# Capitalized word is followed by another
```

THE PAGE TITLE THE SELLER TYPED → WHAT IT GRABS

```
2011 Honda Accord LX ...      → Honda Accord ✓
2013 Hyunda Sonata GLS ...    → Hyunda Sonata (typo)
2013 hyundai elantra ...      → Contact Information
2007 Honda odyssey lx - Vernon Rockville → Vernon Rockville

# the reliable spot, near the photos:
'2007'
'Honda odyssey lx'
```

Regex does exactly what you say, not what you mean. Lowercase title? The first two Capitalized words are a button, or the town.

Build it one piece at a time

piece	means	piece	means
<code>^</code> <code>\$</code> + <code>re.MULTILINE</code>	start / end of a line	<code>\n</code>	the line break
<code>\d{2}</code>	exactly two digits	<code>\s+</code>	one or more spaces
<code>(?:19 20)</code>	19 or 20, grouped, not captured	<code>(...)</code>	capture group → <code>.group(1)</code>

```
step 1  a line that is only a year      r"^(?:19|20)\d{2}$"          → matches '1980'
step 2  ...followed by the next line    r"^(?:19|20)\d{2}\n"
step 3  capture the first word (make)    r"^(?:19|20)\d{2}\n([A-Za-z-]+)"
step 4  a space, then capture the model  your turn
```

Why the year line is safe: Craigslist's form writes it, so it is always 4 digits (never '23). Quirk: now and then the make/model line is one word, like **Pontiac** (no model).

Test fast at regex101.com (Python flavor). Paste 20 lines of an ad as the test string and watch the match light up.

Fix it three ways, then re-score

WAY 1 GIVEN · WAY 2 IS YOURS · WAY 3 STRETCH

```
# Way 1: plain Python (given)
if len(line) == 4 and line.isdigit():
    words = next_line.lower().split()

# Way 2: the same idea in ONE regex line
BETTER_MAKE_MODEL_RE = re.compile(
    r"YOUR PATTERN HERE", re.MULTILINE)

# Way 3: fix typos with a list of makes
difflib.get_close_matches("hyunda", KNOWN_MAKES)
# ['hyundai']
```

THE NOTEBOOK RE-SCORES YOUR 5 CARS


field	original	way 1	way 2	way 3
make	0.2	0.8	?	1.0
model	0.4	1.0	?	1.0

Dr. Wanik's Hartford sample, 5 cars

Stretch: run it on **all** your cars. How many ads does your pattern miss, and why?

Same idea, three tools. Rules are fast and cheap, and brittle. In A07, Gemini tries to beat way 3.

Ship your regex: **branch** → **PR** → **deploy**



- Only **main** deploys: the workflows run on main, and Google Cloud only trusts main.
- Keep your own analysis in **separate files** (notebooks) so **Sync fork** still works when Dr. Wanik updates the course code.

That's CI/CD: code change → review → automatic deploy. No clicking around in the cloud.

Good advice

- **Colab pop-up:** pick the Gmail you used for **Google Cloud**, not necessarily your UConn account.
- **listings_master.csv not found?** Force run the extractor, then materialize-master; wait for **Success**.
- **Your regex matches nothing?** Print the lines with `repr()`: hidden spaces and line breaks show up.
- **Only 10 cars?** That's one scrape. It adds up to 10 **new** cars every hour; leave it running.
- **Pipeline not ready?** `USE_BACKUP_DATA = True` loads Dr. Wanik's Hartford cars. Do the lab now, fix your pipeline after.

DEFINITION OF DONE

Submit Lab 3 on HuskyCT

- ☐ Part 2: search results + 2–3 sentences on what the raw text keeps and loses
- ☐ Part 3: your **truth table** for 5 cars + the **before** accuracy table
- ☐ Part 4: your **BETTER_MAKE_MODEL_RE** line + the **before vs after** table
- ☐ Part 5: truth table **saved for A07** (Gemini gets scored on the same 5 cars) + three short answers
- ☐ File → Download → **.ipynb** → **Submit Lab 3** (due **Thu Oct 8, 11:59 PM**)

Leave your pipeline running. Next: A06 (update the ETL), then A07, where Gemini reads the ads instead of a regex.