

From a Craigslist ad to a price prediction, every hour

Your pipeline runs on Google Cloud while you sleep. Nobody presses a button.

Real data from Dr. Wanik's Hartford cars
(the prediction panel is an illustration)



Everything lands in **your own bucket** on Google Cloud: raw text, JSON, the master table, predictions. Nothing is lost when an ad disappears.

Change code, and the cloud updates itself: branch → pull request → merge → GitHub Actions redeploys. That is CI/CD.

Now it is yours. What will you build on top?

Fix the extractor

LAB 3 · A06

One line of regex turns "Contact Information" into real makes and models, and you ship it with a pull request.

Mine more features

100 REGEX IDEAS

Condition, title status, cylinders, photos, one owner, new tires: all sitting in the ad, waiting for a pattern.

Regex vs. Gemini

A07

Score the LLM on your own truth table. Where is it smarter? Where is plain regex faster, cheaper, and just as right?

Clean before you train

LAB 3 · MIDTERM

Flag the suspicious rows (unknown make, 99,999 miles) and keep them out of your model.

Explain every price

XAI · MODULE 2

SHAP shows why the model priced a car the way it did: the mileage, the year, that rebuilt title.

Go bigger

THE CLASS DATASET

50 of you scrape 50 different areas. Pool them and train one master model on all 50 areas.